

Virtualised filenames – an alternative to docker.



Portable vs non-portable applications

Portable application	Non-portable application
<pre>If (Posix API available or Win API available) { Take user input; Process user input; Produce output; }</pre>	<pre>If (not on x86) goto fail; If (not on linux) goto fail; If (not on linux 3.4+) goto fail; If (libc is not glibc) goto fail; If (glibc is not version 2.21+) goto fail; If (/usr/bin/python not there) goto fail; If (/usr/bin/python --version is not exactly 2.7.10) goto fail; If (current executable's filepath is not in /usr/bin) goto fail; If (you're lucky enough that none of the above) { Take user input; Process user input; Produce output; }</pre>

I want this program installed & running by Monday




most unportable < least unportable

Buy new (possibly old) hardware

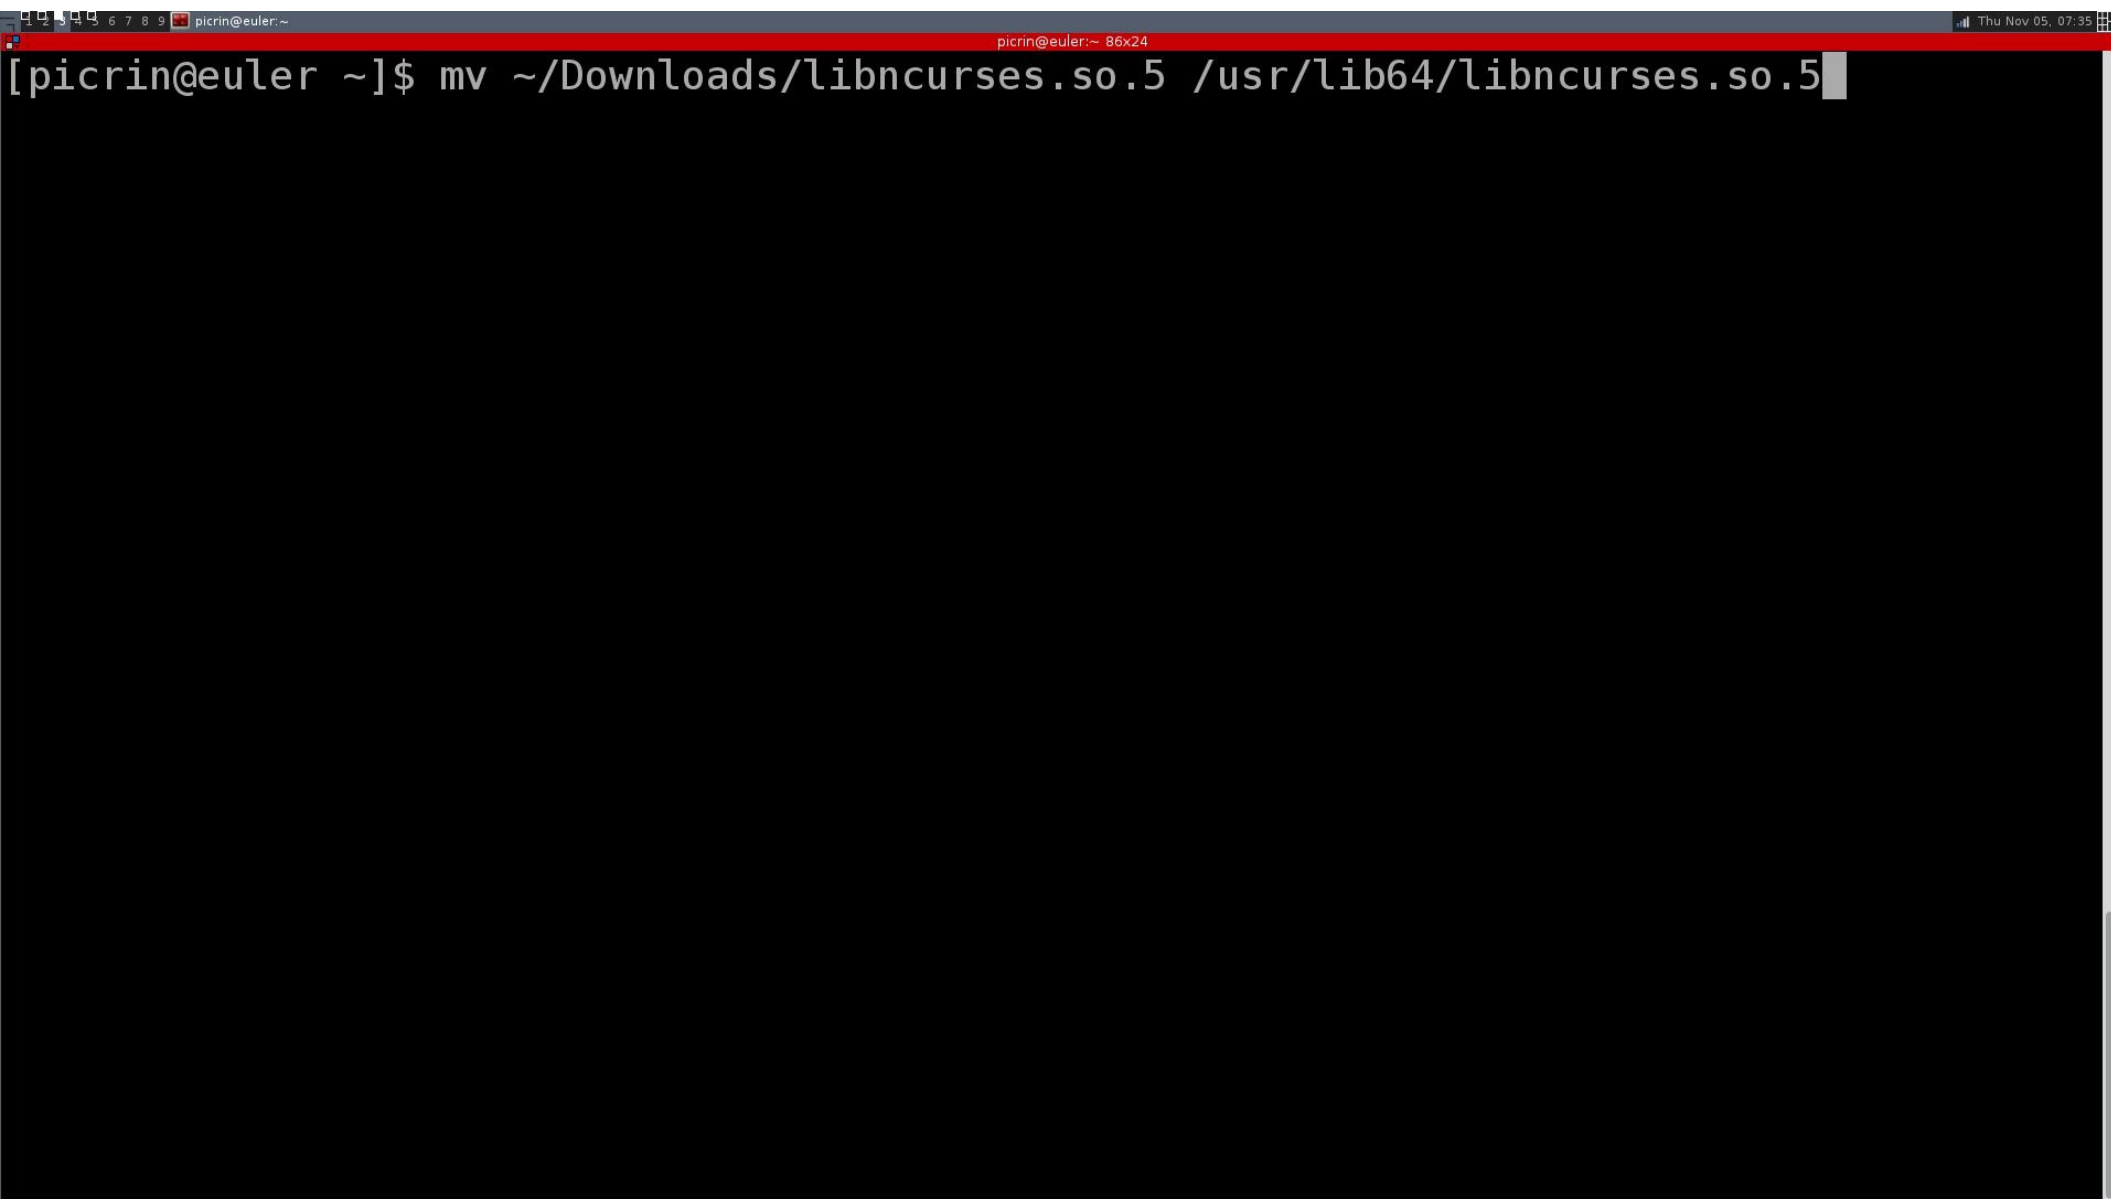
Emulate the architecture

Emulate the OS

Emulate fs, network stack and process tree.




There is one more way...

A terminal window with a red title bar. The title bar contains window control icons on the left, the text 'picrin@euler:~' in the center, and 'Thu Nov 05, 07:35' on the right. The terminal content shows a command prompt '[picrin@euler ~]\$' followed by the command 'mv ~/Downloads/libncurses.so.5 /usr/lib64/libncurses.so.5'. A white cursor is positioned at the end of the command.

```
[picrin@euler ~]$ mv ~/Downloads/libncurses.so.5 /usr/lib64/libncurses.so.5
```


Painful way



Rpm/dpkg to the rescue?

```
picrin@euler:~$ ls /usr/bin/python*
/usr/bin/python          /usr/bin/python2-config  /usr/bin/python3.4m      /usr/bin/python-config
/usr/bin/python2         /usr/bin/python2-coverage /usr/bin/python3-chardetect /usr/bin/python-coverage
/usr/bin/python2.7       /usr/bin/python3         /usr/bin/python3-mako-render
/usr/bin/python2.7-config /usr/bin/python3.4       /usr/bin/python3-pyinotify
picrin@euler ~]$
```

Not Really



Hack into libc a new call vfn

A process makes a call
`vfn("/usr/bin/python",
"/home/picrin/python2_5")`

From now on each call to
`open()`, `stat()`, `exec()`, etc.
will execute as normal, EXCEPT if
path is `/usr/bin/python`, in which
case they'll execute as if path
was `/home/picrin/python2_5`

Each process can execute vfn
multiple times, to virtualise as
much (or as little) of the
filesystem as it wishes.

Each process inherits the vfn
translation table from its parent
process.

A process doesn't make the vfn
call

each call to `open()`, `stat()`,
`exec()`, etc. will execute as
normal.

No process can alter vfn
translations of any other process
(security).

You don't have to be root to alter
your own process's vfn translation
table (sensible).

Nitty-gritty details-pitfalls,
like executing vfn on files with
open descriptors, etc.

vfn translation table possibly has
to be implemented inside kernel

vfn to the rescue!

`rpm -ql python3`

```
/usr/bin/pydoc3  
/usr/bin/pydoc3.4  
/usr/bin/python3  
/usr/bin/python3.4  
/usr/bin/python3.4m  
/usr/bin/pyvenv  
/usr/bin/pyvenv-3.4
```

`rpm -ql python3 | xargs sha256sum`

```
/usr/bin/pydoc3 -> /opt/vfn/8ff8790221fbf907c4c6e2db127664ec59f47e9bfb659252441900dbcc0dccaaf  
/usr/bin/pydoc3.4 -> /opt/vfn/8ff8790221fbf907c4c6e2db127664ec59f47e9bfb659252441900dbcc0dccaaf  
/usr/bin/python3 -> /opt/vfn/4b78808d15be5bebfa7e0d11410d3bca84e370a196ec06d144ed214540b4060f  
/usr/bin/python3.4 -> /opt/vfn/4b78808d15be5bebfa7e0d11410d3bca84e370a196ec06d144ed214540b4060f  
/usr/bin/python3.4m -> /opt/vfn/4b78808d15be5bebfa7e0d11410d3bca84e370a196ec06d144ed214540b4060f  
/usr/bin/pyvenv -> /opt/vfn/9e64cc7d0101933c7288e1866f5a8d70ddef8a121308a4faa987c4c0c00a254c  
/usr/bin/pyvenv-3.4 -> /opt/vfn/9e64cc7d0101933c7288e1866f5a8d70ddef8a121308a4faa987c4c0c00a254c
```

yeah!

Packaging becomes easy.

Dependency reuse
becomes easy.

Software shipping becomes
easy.

Containers become secure (no UID=0).

You no longer need UID=0 to install anything anywhere.

Writing software is still hard ← focus there, portability is
still a big bonus!



wait a minute

Need own kernel module to keep the vfn translation table, work out table inheritance, etc.

Need to hack EVERY libc call, which does a filesystem operation.



Some software STILL compiles against kernel headers, no way to do translation there.

There are multiple libc's to hack (glibc, bionic, bsd, etc.).

What about windows?